

EmpowerLens Relabelling Codebook (TherapistQA)

Sep 27, 2026 · @nayab

Purpose and ground rules

You are relabelling TherapistQA rows that our model flagged as possibly mislabelled. Your job is to decide which cognitive distortion, if any, the text shows. You are not diagnosing anyone.

- **Label independently.** No discussing rows with other annotators at any point before hand-in.
- **You will not see the original labels or the model's prediction.** That is deliberate, so your judgement stays independent.
- **Judge the text, not the person.** Label what the words show. Do not guess at backstory that isn't written.
- **Pace: about 45 seconds per row.** If a row takes more than a minute, set Unsure = Yes, give your best guess, and move on.
- **Read the whole text first, then the distorted part.** The distortion must be supported by the text itself.
- **Keep this codebook open** throughout the session and check the confusable pairs when torn.

Sheet layout

Each annotator gets one workbook. Labels are chosen from dropdowns, not checkboxes, because dropdowns capture primary vs secondary and block typos.

Column	What you do	Allowed values
Row ID	Nothing (read only)	Dataset Id_Number
Text	Read	Patient question text
Distorted part	Read	Span from the original annotation; may be blank
Primary	Required	One of the 10 distortions, No Distortion, or No Distortion (calibrated pessimism)
Secondary	Optional	One of the 10 distortions, or leave blank

Column	What you do	Allowed values
Unsure	Optional	Yes, or leave blank
Notes	Optional	A few words on why you were torn

Validation rules:

- Primary cannot be blank.
- If Primary is either No Distortion option, Secondary must be blank.
- Secondary cannot repeat Primary.
- Distortion values match the dataset's label strings exactly. No Distortion (calibrated pessimism) is our addition; for comparison with the original labels it merges into No Distortion.

A blank distorted part does not mean "no distortion". Judge the text itself.

The 10 distortions

Use the label names exactly as below; they follow the TherapistQA label set. Check the casing against the dataset before building the dropdowns.

Label	The thought does this	Typical cues	Example
All-or-nothing thinking	Sees things in two extremes, with no middle ground	always/never about outcomes, perfect/ruined, total failure	"If I don't get an A, the whole year is wasted."
Overgeneralization	Draws a sweeping rule from one or a few events	every time, nothing ever, everyone	"My friend cancelled once, so nobody wants to spend time with me."
Mental filter	Dwells on one negative detail and ignores the positives (includes discounting good things)	but, only, doesn't count, yes but	"She praised my work, but she noticed one typo, so it went badly."
Should statements	Rigid rules for self or others; guilt or anger when broken	should, must, have to, ought to	"I should be able to handle this without help."

Label	The thought does this	Typical cues	Example
Labeling	Attaches a fixed global label to self or others	I'm a/an..., he's such a...	"I'm a loser."
Personalization	Takes responsibility or blame for things outside one's control	my fault, because of me	"My parents argue because I'm a bad son."
Magnification	Blows up how bad something is, or shrinks the good; catastrophizing	disaster, unbearable, the worst, can't cope	"If I stutter in the meeting, it will be a catastrophe."
Emotional Reasoning	Treats a feeling as proof of a fact	I feel ____, so it must be ____	"I feel stupid, so I must be stupid."
Mind Reading	Assumes what others think, without evidence	they think, she must feel, he probably believes	"My boss thinks I'm incompetent."
Fortune-telling	Predicts a negative future as if it were certain	will, going to, never going to	"I'll never find a job."

Confusable pairs

Most disagreements come from these pairs. Ask the deciding question.

Pair	Deciding question	Pick the first if...	Pick the second if...
Fortune-telling vs Mind Reading	Is it about the future, or about someone's thoughts now?	A future event is predicted	Another person's current opinion is assumed
Fortune-telling vs Magnification	Is it predicting, or inflating size?	"It will happen" is the core	"It would be terrible" is the core
Overgeneralization vs All-or-nothing thinking	Is it a rule from past events, or a two-option view?	Built from one or few past events	Only two outcomes exist, no middle

Pair	Deciding question	Pick the first if...	Pick the second if...
Overgeneralization vs Labeling	Is the conclusion about events or identity?	A pattern of events ("this always happens")	A fixed identity ("I'm a failure")
Mental filter vs Magnification	Is good being ignored, or bad being inflated?	Positives are present but dismissed	One bad thing is blown up, no positives mentioned
Personalization vs Labeling	Is it blame for an event, or a global identity?	"It happened because of me"	"I am a ___"
Emotional Reasoning vs Labeling	Is the feeling the evidence?	"I feel X, so X is true"	The label stands without a feeling as proof
Should statements vs All-or-nothing thinking	Is it a rule, or an extreme outcome?	must/should rule about behaviour	Outcome judged as total success or total failure

When to choose No Distortion

Choose No Distortion when the text is negative or worried but the reasoning holds up. Sadness, fear or frustration alone is not a distortion.

- **Describing facts or feelings.** "My mother died last month and I can't stop crying." Painful, but no reasoning error. Label: No Distortion.
- **Asking for help or information.** "How do I talk to my son about his grades?" Label: No Distortion.
- **Reported speech.** Someone else's distortion quoted by the writer ("my husband says I'm useless") counts only if the writer adopts it. If not adopted: No Distortion.

No Distortion (calibrated pessimism) is for a negative expectation or self-assessment that fits the evidence in the text. Use it when the thought *looks* like Fortune-telling, Magnification or Overgeneralization but the reasoning holds.

- **Grounded in stated evidence.** "I failed the last two driving tests, so I'm nervous about the third."
- **Hedged, not certain.** "I might not get the job" is calibrated. "I'll never get a job" is Fortune-telling.

- **Proportionate.** "Losing this client would hurt our income" is calibrated. "Losing this client would ruin everything" is Magnification.

If the text is negative but makes no prediction or judgement at all, use plain No Distortion, not calibrated pessimism.

Test: could a reasonable, well-informed person think this given what's written? If yes, pick one of the two No Distortion options.

Primary vs secondary

Primary is the distortion that drives the text; secondary is a clearly separate second one. Most rows need only a primary.

1. **Primary = the one the writer's distress hangs on.** If you removed it, would the text lose its main problem? Then it's primary.
2. **Tie-break:** if two seem equal, pick the one in the distorted part as primary.
3. **Secondary only if clearly present.** You should be able to point to different words for it. "Could arguably be" is not enough.
4. **No third label.** If you see three, keep the two strongest and note the third in Notes.
5. **Don't double-count one phrase.** "I'm a failure and always will be" is Labeling plus Fortune-telling only because two separate claims are made.

Session procedure

The session runs 4 hours: a shared overlap set for agreement, then a split main pass. Read this codebook before arriving.

Time	Step	What happens
0:00–0:45	Overlap set	Everyone labels the same 60 rows (the first 60 in every workbook), independently. Used for Fleiss' kappa.
0:45–0:55	Break	
0:55–3:45	Main pass	Each annotator labels their own remaining rows, independently.
3:45–4:00	Hand-in	Save and submit workbooks; list rows marked Unsure.

- Do not change overlap-set labels after the break, even if you'd now label differently.

- Unsure rows go to the psychology annotators for a later adjudication round.
- Unfinished rows are fine; submit what's done rather than rushing the last ones.